

# **RavDevOps Engineering Code Culture & Safety Standard**

Establishing a Google- and NASA-Inspired Software Engineering Practice

Organization: RavDevOps ([ravdevops.com](http://ravdevops.com))

Version: 1.0

Status: Draft for Internal Adoption

This whitepaper defines the engineering culture, coding rules, and operational practices at RavDevOps. It combines proven practices from large-scale environments such as Google with safety-oriented standards inspired by NASA's software engineering guidelines. The objective is to create systems that are reliable, maintainable, observable, and safe by default.

# Table of Contents

<b>Placeholder for table of contents</b>	<b>0</b>
--	----------

# 1. Purpose & Scope

RavDevOps builds and operates systems that power critical workloads for our customers. Reliability, maintainability, and safety are treated as first-class features. This whitepaper defines the engineering code culture we expect of all contributors: employees, contractors, and partners.

The goals of this document are to:

- Define a clear, opinionated engineering culture for RavDevOps.
- Specify non-negotiable coding, testing, and operational rules for production-facing software.
- Align our practices with proven patterns from Google-scale systems and NASA-inspired software safety standards.
- Provide a foundation for automated enforcement by tests, CI/CD, and AI agents.

These rules apply to all production-facing software and infrastructure-as-code owned or operated by RavDevOps. Exceptions must be explicitly documented, justified, reviewed, and time-bounded.

## **2. Core Engineering Principles**

The following principles guide every design, implementation, and operational decision at RavDevOps. They represent the behaviors we want to make easy and automatic via processes, tools, and AI assistants.

### **2.1 Write Boring, Readable Code**

Code should read like a story. Control flow and data flow must be obvious, with variables declared near first use and minimal surprises for the reader. Clever code is a liability; simple, boring code is an asset. We prioritize clarity and predictability over compactness or novelty.

### **2.2 Design for 10x Scale**

Every new system and feature must be safe at ten times today's projected load: traffic, data volume, and organizational use. We assume: more regions, more users, more teams, and more dependencies than initially planned. Designs consider sharding, caching, partitioning, back pressure, and data distribution from the outset.

### **2.3 Favor Statelessness and Immutability**

State is a liability and must be deliberately managed. Services should be stateless between requests, relying on durable data stores, queues, and caches for persisted state. Within a service, prefer immutable data structures and functional-style flows where reasonable to reduce hidden coupling and side effects.

### **2.4 Design Docs Before Meaningful Implementation**

Non-trivial features and system changes require a design document before substantial implementation work. Design docs must clarify the problem, constraints, alternatives, trade-offs, failure modes, rollout plans, and rollback strategies. Design docs serve both as a decision record and as a learning resource for future engineers.

### **2.5 API Stability Is a Contract**

APIs and schemas are contracts with callers. Backward compatibility is the default. Breaking changes must be versioned, communicated, and supported through a migration period. Silent breaking changes are unacceptable.

### **2.6 Test Behavior, Not Implementation**

Tests should verify externally observable behavior and guarantees, not internal implementation details. Internal refactors should not break tests when behavior remains unchanged. Mocks are used only for true external dependencies, not the system under test itself.

## **2.7 Deterministic Builds and Environments**

Builds and tests must be reproducible. Any engineer or CI job should obtain the same result from the same revision. This requires explicit, pinned dependencies, declarative build configurations, and strict control of environment variability.

## **2.8 Design for Failure as the Default**

We assume that networks partition, machines fail, dependencies degrade, and retries happen at the worst possible moment. We design for graceful degradation, bounded retries, timeouts, and idempotent operations to avoid data corruption and cascading failures.

## **2.9 Optimize for P99, Not P50**

Users and systems experience tail latency, not averages. In distributed and microservice architectures, tail latency compounds across call chains. We define and monitor SLOs in terms of P95/P99 latency and error rates and treat regressions in tail metrics as first-class incidents.

## **2.10 Blameless Postmortems and Psychological Safety**

Incidents are treated as systemic failures, not personal failures. Postmortems are blameless, focused on learning, and result in concrete follow-up actions. Engineers must feel safe to escalate issues quickly, report near-misses, and share candid feedback.

## **2.11 Kindness, Humility, and Collaboration**

Technical excellence is necessary but not sufficient. We expect engineers to behave with kindness, humility, and a bias to help. Mentorship, generous code reviews, and knowledge sharing are integral parts of our culture.

## 3. Concrete Coding Rules

This section translates our principles into explicit, enforceable rules. These rules are language-agnostic where possible and are tightened further for safety-critical or infrastructure-critical components.

### 3.1 Control Flow and Complexity

#### **R■CF■1 – Simple Control Flow Only**

- Do not use constructs such as goto or label-based multi-level breaks/continues.
- Avoid deep nesting: target a maximum of three levels of nested conditionals or loops per function.
- Avoid unbounded recursion in production services. Recursion may be used only when the maximum depth is clearly limited and not on critical hot paths.

#### **R■CF■2 – Bounded Loops**

- All loops must have an obvious and documented upper bound in terms of iterations, time, or data volume.
- Long-running or potentially unbounded loops must be interruptible and emit metrics for observability.

#### **R■CF■3 – Function Size Limits**

- Aim for functions of at most 50–60 lines of effective code (excluding comments and trivial braces).
- Large functions must be decomposed into smaller, well-named helpers that express the steps of the algorithm clearly.

### 3.2 Data and State

#### **R■DS■1 – Minimize Local and Shared State**

- Declare variables in the narrowest scope that satisfies readability.
- Avoid large mutable shared objects; prefer immutable data structures or narrow, well-defined interfaces.

#### **R■DS■2 – Stateless Service Rule**

- Statelessness is the default for business logic services. Any state that must survive a request boundary belongs in a datastore, queue, cache, or explicit state machine.
- Do not hide state in in-memory singletons that span requests in production services.

#### **R■DS■3 – No Ad■Hoc Global Business State**

- Globals may be used only for constants, configuration descriptors, and dependency injection wiring.
- Business state must never rely on mutable global variables.

### 3.3 Memory, Resources, and Allocation

## **R■MR■1 – Avoid Dynamic Allocation in Hot Paths**

- Avoid frequent dynamic allocation in high-QPS or latency-critical paths.
- Where appropriate, use object pools, preallocation, or stack-based allocation to minimize GC pressure and fragmentation.

## **R■MR■2 – Explicit Resource Ownership**

- Every resource (file descriptor, connection, handle, lock) must have a single clear owner and explicit lifetime.
- Use language-appropriate patterns (RAII, defer, using) to ensure resources are always released.

## **3.4 Assertions, Errors, and Contracts**

### **R■AE■1 – Assertions for Invariants**

- Non-trivial functions should contain at least one meaningful assertion of invariants or assumptions.
- Assertions must be side-effect free and must not change program behavior in production other than fail-fast modes where appropriate.

### **R■AE■2 – Check All Non■Void Results**

- Non-void, non-trivial function calls must have their return values checked or deliberately ignored with an explicit comment.
- Silent ignoring of error codes or results is prohibited in production code.

### **R■AE■3 – Validate Inputs at System Boundaries**

- External inputs (user inputs, API payloads, message bus events, configuration) must be validated at boundaries.
- Internal modules may rely on established invariants but should assert them.

## **3.5 API and Schema Evolution**

### **R■AP■1 – Backward Compatibility First**

- Do not change existing field semantics without versioning.
- Do not remove fields or parameters from public or widely-used internal APIs without a clear deprecation plan and migration strategy.

### **R■AP■2 – Versioned Breaking Changes**

- Breaking changes require a new versioned endpoint, RPC method, or schema.
- Old versions must be supported for a defined coexistence period and monitored for remaining usage.

### **R■AP■3 – Serialization Compatibility**

- For structured data formats (e.g., JSON, Protobuf), follow forward- and backward-compatible evolution practices.
- Adding optional or nullable fields is preferred over altering or removing existing fields.
- Changes to types or semantics must go through the design review and versioning process.

## 4. Testing and Quality Rules

Testing enforces our contracts and guarantees. We place strong emphasis on fast, deterministic tests that validate behavior under both normal and failure conditions.

### 4.1 Behavior■Driven Testing and the Pyramid

#### **R■TST■1 – Behavior■Focused Tests**

- Tests describe externally observable behavior and guarantees: what the system does given inputs and environment.
- Avoid over-specifying internal structure or mocking internal methods of the system under test.
- Mocks and fakes are reserved for true external dependencies (databases, queues, third-party APIs).

#### **R■TST■2 – Testing Pyramid Composition**

Each service or component must maintain an appropriate blend of tests:

- Unit tests: high volume, fast, isolated; mandatory for core logic.
- Integration tests: verify interactions with real implementations of dependencies or realistic test doubles.
- End■to■end tests: small set of realistic flows to validate overall behavior in an environment close to production.

The majority of tests should be unit tests, complemented by a targeted set of integration and end■to■end tests.

#### **R■TST■3 – Test Speed and Determinism**

- Unit test suites must run quickly enough to be executed frequently during development.
- Tests must be deterministic: avoid dependence on real time, network variability, or external services.
- Where time is involved, use fake clocks or time abstractions to control test behavior.

### 4.2 Static Analysis and Warning Policies

#### **R■SA■1 – Zero Warning Policy for Main Branch**

- Production code must compile or lint with all relevant warnings enabled, at the strictest practical level.
- The main branch must remain free of compiler and linter warnings. New warnings introduced by a change must be resolved before merge.

#### **R■SA■2 – Static Analysis as a Gate**

- At least one strong static analysis tool must be configured for each primary language.
- CI must fail when new static-analysis issues are introduced, unless explicitly waived with a documented justification and time-bound follow-up.

## 5. Reliability and Operations Rules

We treat reliability as a shared responsibility between development and operations. The following rules operationalize the principle that failure is the default and must be managed proactively.

### 5.1 Designing for Failure

#### **R■REL■1 – Bounded Retries with Backoff and Jitter**

- All remote calls that may fail transiently must use bounded retry policies.
- Retries must use exponential backoff with random jitter to avoid thundering herds.
- Retries must respect end-to-end request deadlines and not extend them silently.

#### **R■REL■2 – Idempotent Operations**

- Operations that may be retried must be idempotent or must detect and safely handle duplicate requests.
- Idempotency keys or operation tokens should be used for side-effecting operations when feasible.

#### **R■REL■3 – Timeouts, Circuit Breakers, and Load Shedding**

- Every remote call must define a timeout appropriate to its context.
- Critical dependencies must be protected by circuit breakers or equivalent mechanisms to prevent cascading failures.
- When under stress, services must shed load gracefully, prioritizing core operations and maintaining overall system health.

#### **R■REL■4 – Graceful Degradation**

- When dependencies fail or degrade, services should return partial results, cached responses, or meaningful error messages where possible.
- Designs must explicitly describe degradation behavior for critical dependencies and user-visible features.

### 5.2 SLOs, Latency, and Observability

#### **R■SLO■1 – Service■Level Objectives for Every Service**

- Each production service must define explicit SLOs for latency (P95/P99) and error rate.
- SLOs must be documented in the service design and visible to all stakeholders.

#### **R■SLO■2 – Tail■Focused Monitoring**

- Dashboards must highlight P95/P99 latency, not only averages.
- Alerts should trigger on breaches or sustained regressions in tail latency and error rates.
- Observability (metrics, logs, traces) is mandatory for diagnosing and resolving incidents quickly.

## 6. Culture, Postmortems, and Governance

Tools and rules are only effective when embedded in a healthy engineering culture. This section defines how we learn from failure, manage risk, and govern exceptions.

### 6.1 Blameless Postmortems

#### **R■HUM■1 – Postmortems for Significant Incidents**

- All Sev■1 and Sev■2 incidents must have a written postmortem.
- Postmortems must include a clear timeline, impact assessment, root cause analysis, contributing factors, and action items.
- The focus is on improving systems and processes, not assigning blame to individuals.

#### **R■HUM■2 – Psychological Safety and Escalation**

- Engineers are encouraged and expected to escalate issues quickly when they suspect user impact or systemic risk.
- Reporting near■misses and uncomfortable truths is valued and protected.
- Leaders model blameless behavior and prioritize learning outcomes from incidents.

### 6.2 Governance and Exceptions

Not all rules can apply equally in all contexts. However, exceptions must be explicit and controlled.

#### **R■GOV■1 – Documented Exceptions**

- Any deviation from the rules in this whitepaper must be documented with:
  - Specific rule(s) being waived.
  - Justification for the exception.
  - Additional risk mitigations.
  - Time-bound review date.
- Exceptions must be visible in code review or design documents and approved by a designated senior engineer.

#### **R■GOV■2 – Periodic Review**

- This whitepaper and any long-lived exceptions must be reviewed periodically.
- We expect the rules to evolve as we gain experience, adopt new technologies, and integrate stronger automation and AI assistance.

## **7. Automated Enforcement and Agentic AI**

RavDevOps intends not only to define engineering culture but to encode it into automated tooling and AI agents that guide and enforce these practices throughout the software lifecycle.

### **7.1 Policy-as-Code and CI/CD Gates**

We implement policy-as-code for coding rules, testing expectations, and infrastructure configurations. Continuous integration and delivery pipelines enforce these policies as merge and deploy gates. Examples include:

- Linters and formatters configured with RavDevOps rule sets.
- Static analysis tools for language-specific and cross-cutting rules.
- Infrastructure policy checks for Kubernetes, Terraform, and cloud resources.
- Test coverage, flakiness detection, and performance regression checks.

### **7.2 Agentic AI for Design, Code Review, and Reliability**

We employ AI agents that understand this whitepaper and act as tireless, consistent reviewers and advisors:

- Design Doc Review Agent: checks for clarity, completeness, failure-mode analysis, and alignment with our principles.
- Code Review Agent: reviews pull requests for readability, data-flow clarity, safety rules, and consistency with design docs.
- Testing & Reliability Agent: evaluates test suites for behavior-focus, coverage of failure modes, and deterministic execution.
- Postmortem Assistant: helps extract timelines, contributing factors, and recommended actions from incident data.

### **7.3 Feedback Loop from Production**

AI agents and automated systems consume telemetry from production—metrics, logs, traces, and incidents—to detect patterns of risk and non-compliance. These insights feed back into our design practices, code rules, and tests. Over time, this creates a virtuous cycle where the culture is continuously reinforced and refined by real-world experience.

### **7.4 Alignment with [ravdevops.com](http://ravdevops.com)**

This whitepaper represents the public-facing articulation of RavDevOps engineering culture for [ravdevops.com](http://ravdevops.com). It is intended to inform our customers, partners, and prospective team members about how we build and operate systems. Internally, it serves as a foundation for training, onboarding, performance expectations, and the behavior of our agentic AI systems.

Questions, suggestions, or proposed changes to this whitepaper should be raised through the designated engineering governance process.